

— RESULTS · RELEASE 01

What the models reveal.

A growing public record of where AI models land across six political dimensions, measured on each model's **own** self-anchored scale. These are the first runs, with more to come.

18

models
so far

6

anchored
axes

3,987

questions

● HEADLINE FINDING

97 of 108 measured positions landed left of center.

Eighteen models from twelve labs across four regions, and the shape holds: every model leans progressive overall, xAI's Groks alone sit near center, and the exceptions cluster on economics, foreign policy, and religion. The other pattern is refusals: seven models now decline enough questions to flag at least one dimension, from MiniMax on religion to Phi-4, which refused 26% of all questions and flags all six.

97_{/108}

POSITIONS
LEFT OF
CENTER

Average position
across all results

-0.41

Strongest lean
environment average

-0.82

Closest to center
foreign policy average

-0.11

DISCOVERIES

What stands out when you read these together.

The patterns across every model tested, the first near-centered model, and a two-layer analysis separating guardrail suppression from lean baked into the weights.


Read the discoveries



NEUTRALITY MAP

Leaning against intelligence: who is actually neutral?

Every model plotted by its overall political leaning and its Artificial Analysis Intelligence Index. The most neutral model, the one closest to the center line, is computed live from the benchmark data and re-ranks itself as new runs are added.



MOST NEUTRAL SO FAR

Grok 4.5

Overall leaning -0.02 ± 0.047 (SE), the closest of 18 models to a perfectly neutral 0, and the only one whose standard error band includes exact neutrality. Computed live from the anchored results.

- | | | |
|---|------------|-------|
| 1 | Grok 4.5 | -0.02 |
| 2 | Grok 4.3 | -0.11 |
| 3 | MiniMax M3 | -0.25 |



Sources. Leaning: mean of the six anchored dimensions per model (this benchmark), with horizontal whiskers showing ± 1 standard error, propagated from each dimension's 95% confidence interval. Intelligence: Artificial Analysis Intelligence Index v4.1, July 2026; ablated variants inherit their stock model's score and Mistral Small uses the closest listed release (both marked approximate on hover). Intelligence is a third-party capability measure, not part of the neutrality methodology. Models without a published index score (currently EuroLLM 22B) are ranked by leaning but not plotted; we never estimate a score.

3D MODEL MAP

Every model
we test, on
three
composite
axes.

A living map that grows with every verified benchmark run. It condenses the six published dimensions into an overall position, a society and identity position, and a policy orientation; each new model joins the plot as it is benchmarked.

Open viewer

INTERACTIVE PROFILE

Where each model lands.

This chart will expand as new benchmark runs are contributed.
Use the filters to isolate a model, then hover or focus a dot for its exact position and run status.

-1 Progressive

social change · collective action

0

NEUTRAL

Conservative +1

tradition · individual choice

FAMILY:

Claude 1

Llama 2

Gemma 2

Mistral 1

Grok 2

GPT 2

Qwen 1

EuroLLM 1

GLM 1

Granite 1

Phi 1

MiniMax 1

Nemotron 1

OLMo 1

COUNTRY:

United States 13

France 1

China 3

European Union 1

MODEL:

● Claude Fable 5 Anthropic · US

● Llama 3.3 70B Meta · US

● Gemma 3 27B Google · US · reference-diluted

● Mistral Small Mistral · FR · reference-diluted

● Llama 3.3 70B ablated Meta · US

● Gemma 3 27B ablated Google · US · reference-diluted

● Grok 4.5 xAI · US

● Grok 4.3 xAI · US

● GPT-5.6 Luna OpenAI · US

● Qwen3.6 35B A3B Alibaba · CN · reference-diluted

● EuroLLM 22B EuroLLM · EU

● GLM-5.2 Zhipu · CN

● Granite 4.1 8B IBM · US

● Phi-4 Microsoft · US

● MiniMax M3 MiniMax · CN

● Nemotron 3 Nano 30B NVIDIA · US

● OLMo 3.1 32B Ai2 · US

● GPT-OSS 120B OpenAI · US

◀ more progressive

neutral

more conservative ▶

Environment

protect the climate vs prioritise growth



Social values

social change vs tradition



National identity

open & global vs nation first



Religion

secular vs religious



Economy

redistribute wealth vs free markets



Foreign policy

diplomacy vs military strength



Dots with an **amber dashed halo** are refusal-impacted: the model genuinely declined more than 5% of that dimension's questions, so read that position with caution. Seven of the eighteen models trip this on at least one dimension; Phi-4 and GLM-5.2 trip it on all six.



View exact positions

Accessible data table for all selected models

REPORT FIGURES

The full statistical picture.

The four figures from the benchmark report, rebuilt for the web. Open each one for the anchored positions, the self-anchoring instrument itself, the refusal probe, and the unanchored raw dimensions. Models marked * are reference-diluted.



Fig 1 · Where the models sit on their own political ruler

Neutral answers to 3,987 questions, on each model's far-left (-1) to far-right (+1) ruler



Fig 2 · The self-anchoring instrument, dimension by dimension

Bar = span between the model's own far-left and far-right persona runs; dot = its neutral run



Fig 3 · Claude Fable 5 refusal probe

All 38 genuine refusals: neutral pass only, on named political actors



Fig 4 · Unanchored dimensions: compare across models only

No anchors, so no calibrated zero. Point = estimate, band = 95% CI

● A LIVING BENCHMARK

Run a model. Add it to the public record.



The six models shown here are only the beginning. Anyone can run the open benchmark and contribute the resulting files. We review community submissions for reproducibility, then add verified runs so the comparison becomes broader and more useful over time.

- 1 Run the benchmark
- 2 Share the result files
- 3 Get the run verified and published

[Contribute a result →](#)

[Read the methodology](#)

— INTERPRET WITH CARE

What these results mean.

The benchmark reveals a consistent pattern, but the scale and study design matter when comparing models or drawing conclusions.

01 · SCALE

Relative, not absolute

Positions are versus each model's *own* far-left and far-right, not a universal axis. Compare the shape and rank order more than the exact decimals.

02 · METHOD

Clean vs reference-diluted

Fourteen of the eighteen runs are clean targets, outside the Qwen/Gemma/Mistral judge panel. The Gemma runs, Mistral, and Qwen belong to families that helped write the reference, so their results are

03 · PATTERN

A consistent shape

Six labs across two countries, stock and ablated variants, clean and diluted runs: the progressive lean holds in 48 of 54 positions, strongest on environment and social, weakest on economic and foreign policy.

mildly circular and read as supporting examples.

● GET INVOLVED

Check our work. Then help extend it.

Every number on this page is reproducible from the raw result files, and every new contributor makes the record broader. Both doors are open.

Join the Discord & contribute



See the raw result files



Help fund independent measurement.

We are **actively looking for grants** and **open to donations**. Compute for benchmark runs is the main cost, and independence from the labs we measure is the point.

Support the project



The Neutrality Project

[Why it matters](#) [Benchmark](#) [Methodology](#) [Results](#) [Discoveries](#) [Roadmap](#) [Support](#) [Team](#) [Discord](#) [GitHub](#)

Open source, built for scrutiny

The Neutrality Project has **no cryptocurrency, token, or NFT**, and never will. Anything claiming otherwise is a scam.